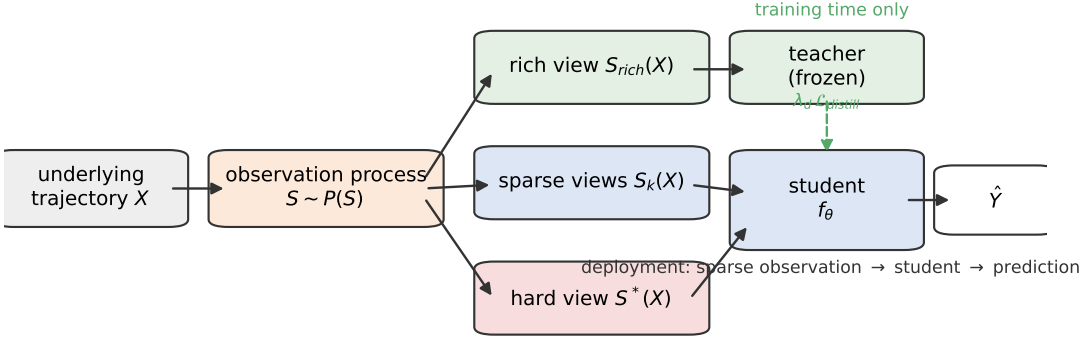
training time only
underlying trajectory X
observation process $S \sim P(S)$
rich view $S_{rich}(X)$
sparse views $S_k(X)$
hard view $S^*(X)$
teacher (frozen)
$\lambda_d \mathcal{L}_{\text{pristill}}$
student f_θ
$\hat{Y}$
deployment: sparse observation → student → prediction
hard mining: search over valid S for the worst case